

Dante Mazza

425-542-4364 · dantemazzal@gmail.com · dantemazza.com
github.com/dantemazza · linkedin.com/in/dante-mazza

SKILLS

Languages: Python, Java, Scala, Kotlin, C, C++, SQL

Frameworks & Data: Spring Boot, Flask, Spark, Hadoop, Delta Lake, Iceberg, Redis, Kafka

Infrastructure: AWS (EMR, EC2, RDS, Lambda, Step Functions), Docker, Kubernetes, Terraform, Linux

Tools: CI/CD (GitLab, GH Actions), Prometheus/Grafana, CloudWatch, Splunk, Claude Code

EXPERIENCE

Splunk

Toronto, Canada

Software Engineer II

Oct 2024 – Present

- Developing scalable, performant (50 TB/day) AWS-hosted Kubernetes + Java/Spring microservices with PostgreSQL/RDS, SQS, Prometheus, GitLab CI/CD, and Terraform.
- Technical lead on a 6 month+ rearchitecture of a Spark ML security detection platform. Designed the data lake, ETL pipeline with feature store, and CI/CD build process, and ran performance PoCs on Parquet writer configs, Delta/Iceberg, and feature computation on EMR.
- Led implementation across 3 engineers and 5 partner teams, improved MTTD 90% (30 to 3 hours), and shipped 1 month ahead of schedule. Mentored 5+ junior engineers and interns.

Software Engineer I

Jul 2023 – Oct 2024

- Built and maintained Spark/Scala batch pipelines using Delta Lake on S3, AWS EMR, and Step Functions to process 1PB+ of logs (VPN, Windows, cloud storage), detecting anomalies across 30M+ users and devices daily.
- Designed and implemented an E2E ML model testing framework (Python, AWS, GitLab CI/CD) and advocated for its adoption, cutting testing and deployment time 98%, from weeks to hours.
- Cut team AWS costs 30% (\$150K/year) by optimizing Spark/Scala ETL pipelines (Glue, Kinesis, Delta Lake) and downstream Spark ML jobs on EMR through performance tuning and code optimization.

Software Engineer Intern

Jan 2022 – Apr 2022

- Prototyped multi-tenant Spark job execution with AWS (EMR, IAM, STS, Step Functions, Lambda), Terraform, Python, Docker, and Java/Spring. Automated tenant provisioning with shell + GitLab/internal APIs, saving 6 hours of onboarding per customer.

Amazon Web Services (Keyspaces)

Bellevue, WA

Software Development Engineer Intern

Aug 2022 – Dec 2022

- Authored a transactional query service in Kotlin using internal AWS services for a managed Cassandra product. Researched open-source Cassandra batch-query functionality and designed for parity.

Wisedocs

Toronto, Canada

Machine Learning Engineer

May 2020 – Sep 2021

- Implemented and optimized ML models (CNNs with TensorFlow/Keras, NLP with Hugging Face, k-means with scikit-learn) for a 2x accuracy improvement. Built Python document clustering that cut prediction time 85%, and a pre-serialized Keras batch-training method that cut training time 90%.

PROJECTS

Shipsolver

shipsolver.com

Sole maintainer of a freight-forwarding B2B suite

Oct 2025 – Present

- Building and operating a multi-tenant B2B freight platform processing 5,000+ deliveries into invoices monthly for 2 active customers. Stack: Python/FastAPI, Next.js/React, AWS (RDS, EC2).

Tricks & Trumps

tricksandtrumps.com

Real-time multiplayer rubber bridge

Jan 2024 – Present

- Engineered a real-time 4-player Bridge app (Flask/Socket.io, React): complete rules engine and Redis-backed state machine.

EDUCATION

University of Waterloo, BAsC in Computer Engineering

Sep 2018 – May 2023

Dean's Honour List (first year). Relevant coursework: Distributed Systems, Compilers, Operating Systems, Networks.